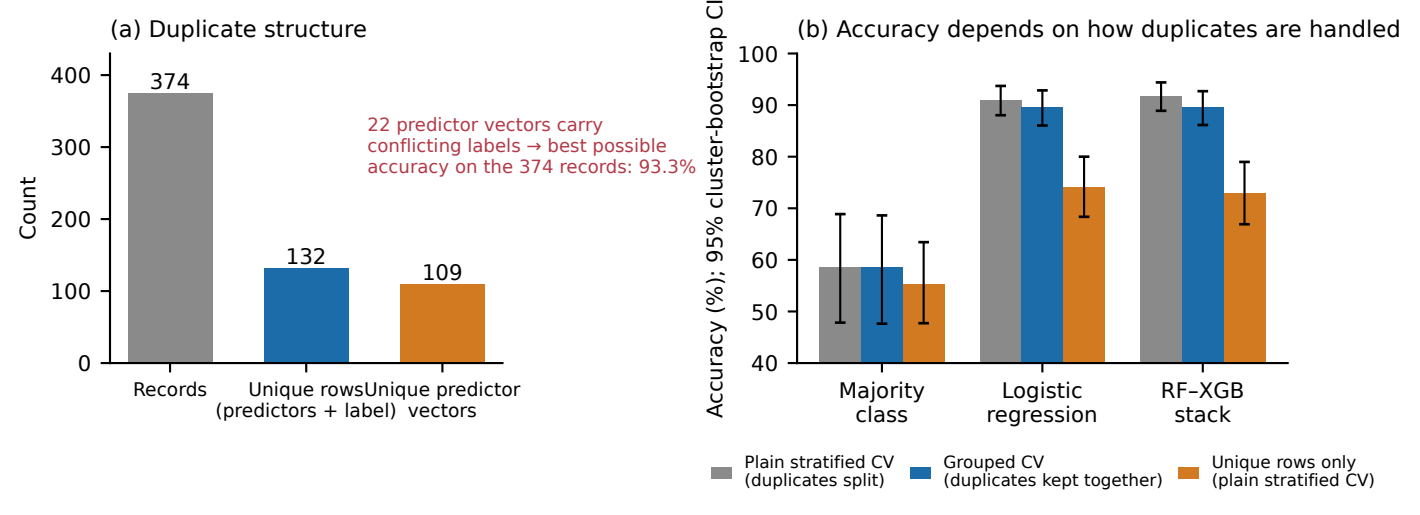
(a) Duplicate structure
Count
400
300
200
100
0
374
132
109
22 predictor vectors carry
conflicting labels → best possible
accuracy on the 374 records: 93.3%
Records
Unique rows
(predictors + label)
Unique predictor
vectors
(b) Accuracy depends on how duplicates are handled
Accuracy (%); 95% cluster-bootstrap CI
100
90
80
70
60
50
40
Majority
class
Logistic
regression
RF–XGB
stack
Plain stratified CV
(duplicates split)
Grouped CV
(duplicates kept together)
Unique rows only
(plain stratified CV)